



AirCopBench: A Benchmark for Multi-drone Collaborative Embodied Perception and Reasoning



Jirong Zha¹, Yuxuan Fan², Tianyu Zhang³, Gen Chen⁴, Yingfeng Chen⁵, Chen Gao⁶, Xinlei Chen¹

The 40th Annual AAAI Conference on Artificial Intelligence
JANUARY 20 - JANUARY 27, 2026 | SINGAPORE

¹Shenzhen International Graduate School, Tsinghua University
²The Hong Kong University of Science and Technology (Guang Zhou)
³School of Electrical and Electronic Engineering, Nanyang Technological University
⁴College of Software, Jilin University
⁵Weiyang College, Tsinghua University
⁶BNRist, Tsinghua University



Project Page



Github



Appendix





Homepage

1. Introduction: MLLM for Collaborative Perception and Reasoning

➤ **Motivation.** Although *multi-drone systems* provide superior coverage, robustness, and coordination, existing benchmarks largely *evaluate MLLMs* only in single-agent vision settings. Most *multi-image benchmarks* focus on perception with high-quality images, failing to capture the complexity of *egocentric collaboration* and *real-world degraded conditions*.

➤ **Contribution.** The *first benchmark* for semantic embodied aerial collaborative perception and reasoning.



2. Task Setting: 4 Dimensions and 14 Types

Scene Understanding

1. Scene Description
Q: What is the dominant environmental feature in UAV1's view?
A: A urban road with a green grassy area.

2. Scene Comparison
Q: Which UAV perspective has more vehicles? A: UAV2.
5 cars, 13 cars

3. Observing Posture
Q: Which perspective is closer to the ground?
A: UAV1 is closer with a clearer view.

Object Understanding

4. Object Recognition
Q: How many types of objects are there? A: 2 types.

5. Object Counting
Q: How many vehicles are visible on the road? A: 5 cars.

6. Object Grounding
Q: Where is the white van?
A: The white van is in a different lane behind the red car.

7. Object Matching
Q: Which object in UAV2 is the person standing near the traffic light in UAV1? A: The person in white standing by the red-and-white sidewalk edge in UAV2.

Perception Assessment

8. Quality Assessment
Q: How would you rate the overall image quality?
A: Good with minor blur.

9. Usability Assessment
Q: Is this image usable for object detection tasks?
A: Moderately usable with some occlusions.

10. Causal Assessment
Q: What is the main factor that affects the object detection performance of this image?
A: Partial occlusion by trees.

Collaborative Decision


11. When to Collaborate
Q: Should UAV2 communicate with others for more info now?
A: Yes.

12. What to Collaborate
Q: What info should UAV1 share with UAV2?
A: Details of the red car's position and movement.

13. Who to Collaborate
Q: Which UAV should UAV2 collaborate with?
A: UAV1.

14. Why to Collaborate
Q: Why should UAV2 collaborate with UAV1?
A: To overcome partial occlusion of the red car.

3. Dataset Statistics



- **2.9k+** multi-view images (simulator + real-world + derived data)
- **14.6k+** visual question answering (VQA) pairs (14 sub-tasks)
- **9** perception degradation scenarios (occlusion, shadow, lighting, blur, noise, loss, etc.) and **4** UAV collaboration configurations

5. Error Analysis

Spatial Reasoning Error

Confused "opposite" and "ahead".

Question: What is the relative position of the white car compared to the bus?
A. The white car is driving in the adjacent lane directly **opposite** the bus. **GT**
C. The white car is **ahead of** the bus in the adjacent lane. **MLLM**

Perception Hallucination Error

Question: Which type of vehicle is most frequently observed in the parking lot?
A. **SUVs** **MLLM** Misclassifying object due to perceptual bias in visual features.
B. **Compact cars** **GT**

Multi-image Understanding Error

Mismatching motion states.

Question: Which object in UAV2 corresponds to the dark gray car approaching the intersection from the left in UAV1? **MLLM**
A: The dark gray car now seen **turning right** at the intersection in UAV2. **GT**
C. The dark gray car now **approaching** the intersection from the top in UAV2.

4. Dataset Generation

Data Collection

Simulator data
3- / 5- / 6-UAV group
multi-view collaboration

Real-world data
long distance, occlusion, different brightness, shadow, challenging degradations

Derived data
Noise injection (sensor failure), Partial masking (data loss)

Data Annotation

Event-level labeling
Image quality (Very poor, Poor, Fair, Good, Excellent), Perception usability (Yes, No), Perception degradation (occlusion, shadow), Collaborative analysis (when, what, who, why)

Object-level labeling
Object list, Bounding box, Target attribute

Question Generation

Model-based generation
Divided task, Role-playing, CoT prompt, Few-shot

Rule-based generation
Q: Which UAV perspective shows more vehicles? A: UAV2.

Human-based generation
Q: Which UAV perspective is closest to the drone target? D. Equally close. A.

Quality Control

Standard examination
Scoring criteria: Required content, Format consistency, Answer validity, Question length

Blind filtering
VQAs, Models, A1 = Y1 = ... = Yn?

Human refinement
Ambiguous questions, Invalid options, Incorrect answers

6. Evaluation Results

Performance Analysis We evaluate **40** popular MLLMs. The *best-performing* models achieve only **59.17%** accuracy, **24.38%** below human performance, emphasizing the challenges in multi-view collaborative perception and reasoning. A pronounced *task bias* emerges, with strong performance on scene description and comparison but poor results on usability assessment and collaborative timing, while several *open-source models* rival or surpass GPT-4o.

SFT and Sim-to-Real After *fine-tuning*, the models' performance improved by up to **26.97%**. Moreover, fine-tuned MLLMs trained on simulated data show notable *sim-to-real* improvements in real-world tasks.